

haberman-prediction

January 30, 2024

```
[47]: # Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.metrics import confusion_matrix, \
    accuracy_score, classification_report
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
```

```
[2]: col_names=["patient_age", "year_of_operation", "positive_axillary_nodes", \
    "survival_status"];
df=pd.read_csv('haberman_kaggle.csv',names=col_names)
df.head()
```

```
[2]:
```

	patient_age	year_of_operation	positive_axillary_nodes	survival_status
0	30	64	1	1
1	30	62	3	1
2	30	65	0	1
3	31	59	2	1
4	31	65	4	1

```
[3]: df.columns
```

```
[3]: Index(['patient_age', 'year_of_operation', 'positive_axillary_nodes',
        'survival_status'],
        dtype='object')
```

```
[4]: df.describe().transpose()
```

```
[4]:
```


	count	mean	std	min	25%	50%	75%	\
patient_age	306.0	52.457516	10.803452	30.0	44.0	52.0	60.75	
year_of_operation	306.0	62.852941	3.249405	58.0	60.0	63.0	65.75	
positive_axillary_nodes	306.0	4.026144	7.189654	0.0	0.0	1.0	4.00	
survival_status	306.0	1.264706	0.441899	1.0	1.0	1.0	2.00	

	max
patient_age	83.0
year_of_operation	69.0
positive_axillary_nodes	52.0
survival_status	2.0

```
[5]: x=df[['patient_age', 'year_of_operation', 'positive_axillary_nodes']]
     y=df['survival_status']
```

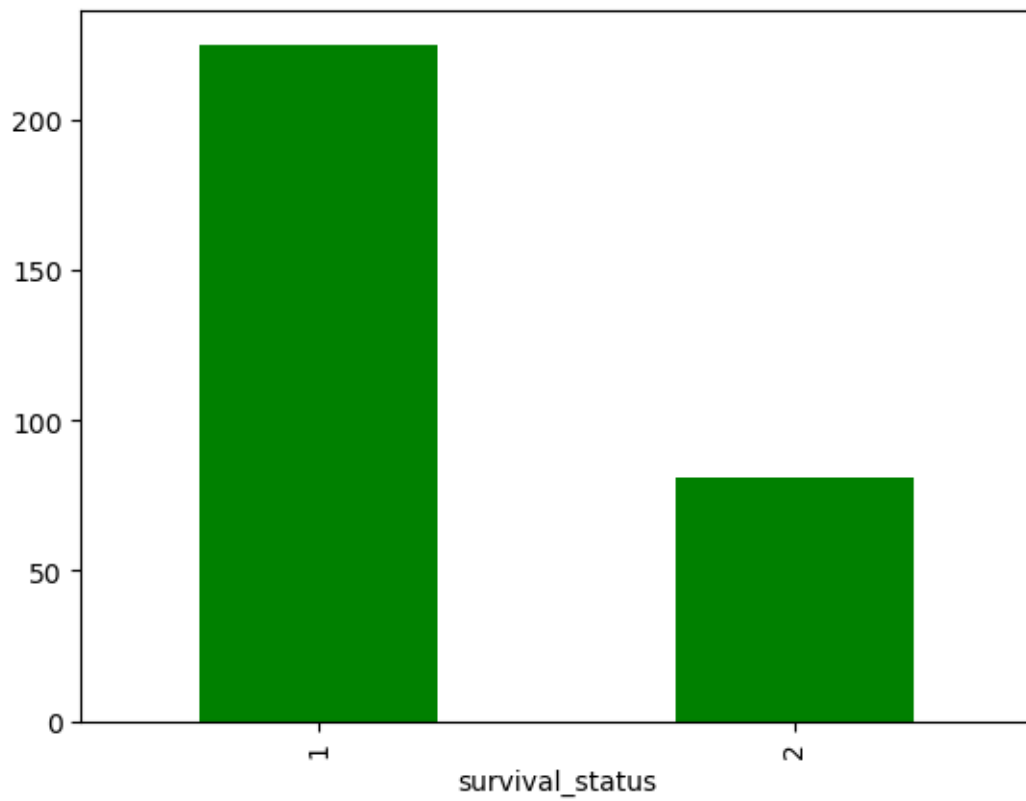
```
[6]: pca = PCA(n_components=2)
     x_pca = pca.fit_transform(x)
```

```
[7]: plt.figure(figsize=(10, 6))
     sns.scatterplot(x=x_pca[:, 0], y=x_pca[:, 1], hue=y, palette='viridis')
     plt.title('Scatter plot after PCA')
     plt.show()
```



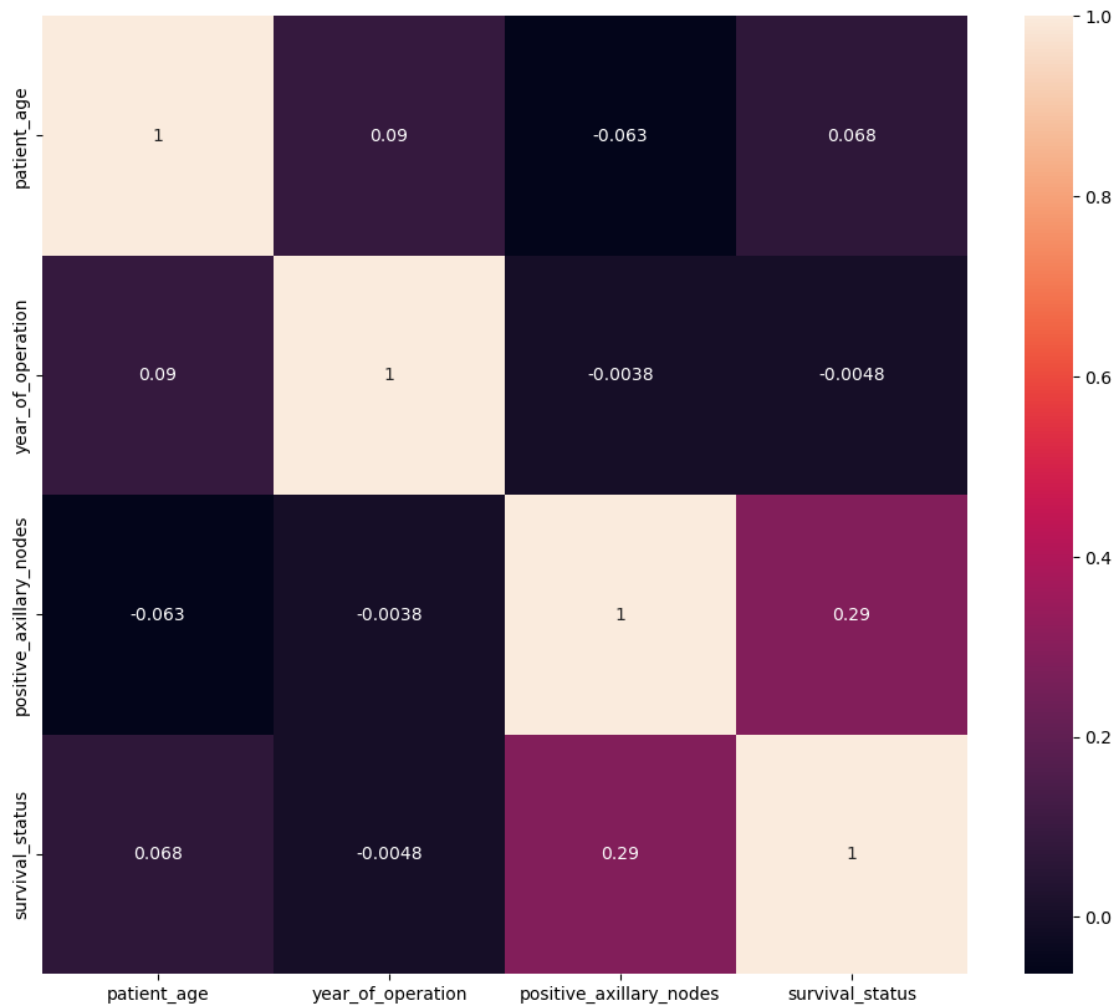
```
[33]: df['survival_status'].value_counts().plot(kind='bar', color='green')
```

```
[33]: <Axes: xlabel='survival_status'>
```



```
[34]: plt.figure(figsize=(12, 10))  
sns.heatmap(df.corr(),annot=True)  
plt.show
```

```
[34]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[10]: x_train,x_test,y_train,y_test,=train_test_split(x,y,test_size=0.
      ↪ 20,random_state=42)
```

```
[11]: scaler=StandardScaler()
      x_train=scaler.fit_transform(x_train)
      x_test=scaler.fit_transform(x_test)
```

0.0.1 KNN

```
[12]: # Hyperparameter tuning
      best_accuracy = 0
      best_k = 1
      for k in range(1, 26):
          neigh = KNeighborsClassifier(n_neighbors=k)
          neigh.fit(x_train, y_train)
          y_pred = neigh.predict(x_test)
```

```
    accuracy = accuracy_score(y_test, y_pred)
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_k = k
best_k
```

[12]: 4

```
[13]: neigh=KNeighborsClassifier(n_neighbors=2, p= 3, weights='uniform')
      neigh.fit(x_train,y_train)
```

[13]: KNeighborsClassifier(n_neighbors=2, p=3)

```
[14]: y_pred=neigh.predict(x_test)
      y_pred=list(y_pred)
      y_test=list(y_test)
```

```
[15]: cm=confusion_matrix(y_test,y_pred)
      cm
```

[15]: array([[39, 5],
 [14, 4]], dtype=int64)

```
[16]: score=accuracy_score(y_test,y_pred)
      score
```

[16]: 0.6935483870967742

```
[17]: neigh=KNeighborsClassifier(n_neighbors=4,weights='uniform')
      neigh.fit(x_train,y_train)
      y_pred_train=neigh.predict(x_train)
      y_pred_test=neigh.predict(x_test)
      from sklearn.metrics import accuracy_score
      score_test=accuracy_score(y_pred_test,y_test)
      print('test accuracy is ',score_test)
      score_train=accuracy_score(y_pred_train,y_train)
      print('train accuracy is ',score_train)
```

test accuracy is 0.7258064516129032
train accuracy is 0.7950819672131147

```
[18]: for i in range(1,26):
      neigh=KNeighborsClassifier(n_neighbors=i,weights='uniform')
      neigh.fit(x_train,y_train)
      y_pred_train=neigh.predict(x_train)
      y_pred_test=neigh.predict(x_test)
      from sklearn.metrics import accuracy_score
```

```
score_test=accuracy_score(y_pred_test,y_test)
print('for',i,'test accuracy is ',score_test)
score_train=accuracy_score(y_pred_train,y_train)
print('for',i,' train accuracy is ',score_train)
```

```
for 1 test accuracy is  0.6451612903225806
for 1  train accuracy is  0.9836065573770492
for 2 test accuracy is  0.6935483870967742
for 2  train accuracy is  0.8237704918032787
for 3 test accuracy is  0.6451612903225806
for 3  train accuracy is  0.819672131147541
for 4 test accuracy is  0.7258064516129032
for 4  train accuracy is  0.7950819672131147
for 5 test accuracy is  0.7096774193548387
for 5  train accuracy is  0.8032786885245902
for 6 test accuracy is  0.6935483870967742
for 6  train accuracy is  0.7786885245901639
for 7 test accuracy is  0.6935483870967742
for 7  train accuracy is  0.7745901639344263
for 8 test accuracy is  0.6774193548387096
for 8  train accuracy is  0.7827868852459017
for 9 test accuracy is  0.6612903225806451
for 9  train accuracy is  0.7663934426229508
for 10 test accuracy is  0.6774193548387096
for 10  train accuracy is  0.7868852459016393
for 11 test accuracy is  0.6935483870967742
for 11  train accuracy is  0.7827868852459017
for 12 test accuracy is  0.7096774193548387
for 12  train accuracy is  0.7663934426229508
for 13 test accuracy is  0.6774193548387096
for 13  train accuracy is  0.7827868852459017
for 14 test accuracy is  0.6774193548387096
for 14  train accuracy is  0.7704918032786885
for 15 test accuracy is  0.6774193548387096
for 15  train accuracy is  0.7827868852459017
for 16 test accuracy is  0.6935483870967742
for 16  train accuracy is  0.7704918032786885
for 17 test accuracy is  0.6774193548387096
for 17  train accuracy is  0.7868852459016393
for 18 test accuracy is  0.6935483870967742
for 18  train accuracy is  0.7786885245901639
for 19 test accuracy is  0.6935483870967742
for 19  train accuracy is  0.7786885245901639
for 20 test accuracy is  0.7096774193548387
for 20  train accuracy is  0.7704918032786885
for 21 test accuracy is  0.6935483870967742
for 21  train accuracy is  0.7745901639344263
```

```

for 22 test accuracy is 0.7096774193548387
for 22 train accuracy is 0.7827868852459017
for 23 test accuracy is 0.7258064516129032
for 23 train accuracy is 0.7827868852459017
for 24 test accuracy is 0.6935483870967742
for 24 train accuracy is 0.7786885245901639
for 25 test accuracy is 0.6935483870967742
for 25 train accuracy is 0.7786885245901639

```

```

[19]: from sklearn.model_selection import GridSearchCV
parameters = {'n_neighbors': [1,2,3,4], 'weights': ['uniform', 'distance'], 'p':
    ↪ [3, 4, 5]}

classifier = KNeighborsClassifier()
clf = GridSearchCV(classifier, parameters, cv=8)

clf.fit(x_train, y_train)
clf.best_params_

```

```

[19]: {'n_neighbors': 2, 'p': 4, 'weights': 'uniform'}

```

```

[20]: best_params = clf.best_params_
classifier_best = KNeighborsClassifier(**best_params)

# Fit the classifier with the best parameters on the training data
classifier_best.fit(x_train, y_train)

# Predictions on test and training sets
y_test_pred = classifier_best.predict(x_test)
y_train_pred = classifier_best.predict(x_train)

# Calculate accuracies
acc_test = accuracy_score(y_test_pred, y_test)
print('Test accuracy is ', acc_test)

acc_train = accuracy_score(y_train_pred, y_train)
print('Train accuracy is ', acc_train)

```

```

Test accuracy is 0.6774193548387096
Train accuracy is 0.8155737704918032

```

0.0.2 Decision Tree

```

[21]: from sklearn.tree import DecisionTreeClassifier

```

```

[22]: dt_classifier = DecisionTreeClassifier()
dt_classifier.fit(x_train, y_train)

```

```
[22]: DecisionTreeClassifier()
```

```
[23]: dt_pred = dt_classifier.predict(x_test)
```

```
[24]: dt_accuracy = accuracy_score(y_test, dt_pred)
      dt_accuracy
```

```
[24]: 0.6612903225806451
```

```
[25]: dt_param_grid = {'criterion': ['gini', 'entropy'], 'max_depth': [2, 5, 10, 15]}
      dt_grid_search = GridSearchCV(dt_classifier, dt_param_grid, cv=5)
      dt_grid_search.fit(x_train, y_train)
```

```
[25]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': [2, 5, 10, 15]})
```

```
[26]: best_params = dt_grid_search.best_estimator_
      best_params
```

```
[26]: DecisionTreeClassifier(max_depth=2)
```

```
[27]: dt_pred_tuned = best_params.predict(x_test)
      dt_accuracy_tuned = accuracy_score(y_test, dt_pred_tuned)
      dt_accuracy_tuned
```

```
[27]: 0.7258064516129032
```

```
[28]: for i in range(1, 26):
      dt_classifier = DecisionTreeClassifier(max_depth=i, random_state=42)
      dt_classifier.fit(x_train, y_train)

      y_pred_train = dt_classifier.predict(x_train)
      score_train = accuracy_score(y_pred_train, y_train)
      print('For max_depth =', i, 'train accuracy is', score_train)

      y_pred_test = dt_classifier.predict(x_test)
      score_test = accuracy_score(y_pred_test, y_test)
      print('For max_depth =', i, 'test accuracy is', score_test)
```

```
For max_depth = 1 train accuracy is 0.7540983606557377
For max_depth = 1 test accuracy is 0.6935483870967742
For max_depth = 2 train accuracy is 0.7909836065573771
For max_depth = 2 test accuracy is 0.7258064516129032
For max_depth = 3 train accuracy is 0.8032786885245902
For max_depth = 3 test accuracy is 0.7258064516129032
For max_depth = 4 train accuracy is 0.8073770491803278
For max_depth = 4 test accuracy is 0.6935483870967742
```



```

For max_depth = 5 train accuracy is 0.8401639344262295
For max_depth = 5 test accuracy is 0.6774193548387096
For max_depth = 6 train accuracy is 0.8565573770491803
For max_depth = 6 test accuracy is 0.6612903225806451
For max_depth = 7 train accuracy is 0.8770491803278688
For max_depth = 7 test accuracy is 0.6774193548387096
For max_depth = 8 train accuracy is 0.9057377049180327
For max_depth = 8 test accuracy is 0.6774193548387096
For max_depth = 9 train accuracy is 0.9221311475409836
For max_depth = 9 test accuracy is 0.6774193548387096
For max_depth = 10 train accuracy is 0.9385245901639344
For max_depth = 10 test accuracy is 0.6129032258064516
For max_depth = 11 train accuracy is 0.9385245901639344
For max_depth = 11 test accuracy is 0.6451612903225806
For max_depth = 12 train accuracy is 0.9631147540983607
For max_depth = 12 test accuracy is 0.6451612903225806
For max_depth = 13 train accuracy is 0.9631147540983607
For max_depth = 13 test accuracy is 0.6451612903225806
For max_depth = 14 train accuracy is 0.9754098360655737
For max_depth = 14 test accuracy is 0.6129032258064516
For max_depth = 15 train accuracy is 0.9836065573770492
For max_depth = 15 test accuracy is 0.6129032258064516
For max_depth = 16 train accuracy is 0.9836065573770492
For max_depth = 16 test accuracy is 0.6129032258064516
For max_depth = 17 train accuracy is 0.9836065573770492
For max_depth = 17 test accuracy is 0.6129032258064516
For max_depth = 18 train accuracy is 0.9836065573770492
For max_depth = 18 test accuracy is 0.6129032258064516
For max_depth = 19 train accuracy is 0.9836065573770492
For max_depth = 19 test accuracy is 0.6129032258064516
For max_depth = 20 train accuracy is 0.9836065573770492
For max_depth = 20 test accuracy is 0.6129032258064516
For max_depth = 21 train accuracy is 0.9836065573770492
For max_depth = 21 test accuracy is 0.6129032258064516
For max_depth = 22 train accuracy is 0.9836065573770492
For max_depth = 22 test accuracy is 0.6129032258064516
For max_depth = 23 train accuracy is 0.9836065573770492
For max_depth = 23 test accuracy is 0.6129032258064516
For max_depth = 24 train accuracy is 0.9836065573770492
For max_depth = 24 test accuracy is 0.6129032258064516
For max_depth = 25 train accuracy is 0.9836065573770492
For max_depth = 25 test accuracy is 0.6129032258064516

```

```

[29]: dt_classifier = DecisionTreeClassifier(max_depth=2)
      dt_classifier.fit(x_train, y_train)

```

```

[29]: DecisionTreeClassifier(max_depth=2)

```

```
[30]: dt_pred = dt_classifier.predict(x_test)
dt_accuracy = accuracy_score(y_test, dt_pred)
dt_accuracy
```

```
[30]: 0.7258064516129032
```

```
[51]: class_rep = classification_report(y_test, y_pred)
class_rep
```

```
[51]: '          precision    recall  f1-score   support\n\n     0.74          0.91          0.82         44\n     18\n\n     accuracy                   0.71         62\n     0.62          0.57          0.56         62\n     62\n\nweighted avg         0.67          0.71          0.67
```

0.0.3 Random forest

```
[31]: from sklearn.ensemble import RandomForestClassifier
```

```
[32]: rf_classifier= RandomForestClassifier(random_state=42)
rf_classifier.fit(x_train, y_train)
```

```
[32]: RandomForestClassifier(random_state=42)
```

```
[35]: y_pred = rf_classifier.predict(x_test)
y_pred
```

```
[35]: array([1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 2, 1, 2, 1,
        1, 1, 1, 1, 1, 1, 2, 1, 1, 1, 1, 2, 1, 1, 2, 1, 1, 1], dtype=int64)
```

```
[36]: accuracy = accuracy_score(y_test, y_pred)
accuracy
```

```
[36]: 0.7096774193548387
```

```
[50]: class_rep = classification_report(y_test, y_pred)
class_rep
```

```
[50]: '          precision    recall  f1-score   support\n\n     0.74          0.91          0.82         44\n     18\n\n     accuracy                   0.71         62\n     0.62          0.57          0.56         62\n     62\n\nweighted avg         0.67          0.71          0.67
```

0.0.4 Logistic regression

```
[52]: from sklearn.linear_model import LogisticRegression
```

```
[53]: logistic_model = LogisticRegression(random_state=42)
logistic_model.fit(x_train, y_train)
```

```
[53]: LogisticRegression(random_state=42)
```

```
[54]: y_pred = logistic_model.predict(x_test)
```

```
[55]: accuracy = accuracy_score(y_test, y_pred)
accuracy
```

```
[55]: 0.7096774193548387
```

0.0.5 Naive Bayes

```
[56]: from sklearn.naive_bayes import GaussianNB
from sklearn.datasets import make_classification
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
```

```
[57]: x, y = make_classification(n_samples=1000, n_features=20, n_classes=2, \
    random_state=42)
```

```
[58]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, \
    random_state=42)
```

```
[59]: naive_bayes_classifier = GaussianNB()
naive_bayes_classifier.fit(X_train, y_train)
```

```
[59]: GaussianNB()
```

```
[60]: y_pred = naive_bayes_classifier.predict(X_test)
```

```
[61]: accuracy = accuracy_score(y_test, y_pred)
accuracy
cm = confusion_matrix(y_test, y_pred)
cm
```

```
[61]: '
      precision    recall  f1-score   support\n\n
0.72      0.90      0.80      0.85         1
107\n\n      accuracy              0.80         200\n
0.81      0.80      0.79      0.79         200\n\n
200\n'
```

```
[62]: class_rep = classification_report(y_test, y_pred)
      class_rep
```

```
[62]: '
      precision    recall  f1-score   support\n\n
0.72      0.90      0.80        0.85         93\n
107\n\n
accuracy              0.80        200\n
0.81      0.80      0.79        0.79        200\n
weighted avg          0.81      0.80      0.79
200\n'
```

```
[ ]:
```